

Payment Fraud Detection Using Machine Learning

1. Importing Libraries, loading dataset

```
In [5]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt
import matplotlib.ticker as mtick
%matplotlib inline
df = pd.read_csv('payment_fraud.csv')
df
```

```
Out[5]:
```

	accountAgeDays	numItems	localTime	paymentMethod	paymentMet
0	29	1	4.745402	paypal	
1	725	1	4.742303	storecredit	
2	845	1	4.921318	creditcard	
3	503	1	4.886641	creditcard	
4	2000	1	5.040929	creditcard	
...
39216	986	1	4.836982	creditcard	
39217	1647	1	4.876771	creditcard	
39218	1591	1	4.742303	creditcard	
39219	237	1	4.921318	creditcard	
39220	272	1	5.040929	paypal	

39221 rows × 6 columns

2. Initial Data Preprocessing

```
In [6]: df.isnull().sum()
```

```
Out[6]: accountAgeDays    0
numItems                0
localTime               0
paymentMethod           0
paymentMethodAgeDays    0
label                  0
dtype: int64
```

```
In [7]: # Check fraud and normal distribution
df['label'].value_counts()
```

```
Out[7]: label
0      38661
1        560
Name: count, dtype: int64
```

3. Outlier Removal

- Using IQR Method

```
In [11]: # removing outliers
# Select only numerical columns
numeric_cols = df.select_dtypes(include='number').columns

# Function to remove outliers using IQR method
def remove_outliers(df, columns):
    for col in columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]
    return df

# Apply the function to your dataframe
df_cleaned = remove_outliers(df, numeric_cols)

# Show the shape before and after to see the data is reduced
print("Original shape:", df.shape)
print("After removing outliers:", df_cleaned.shape)
```


Original shape: (39221, 6)

After removing outliers: (28276, 6)

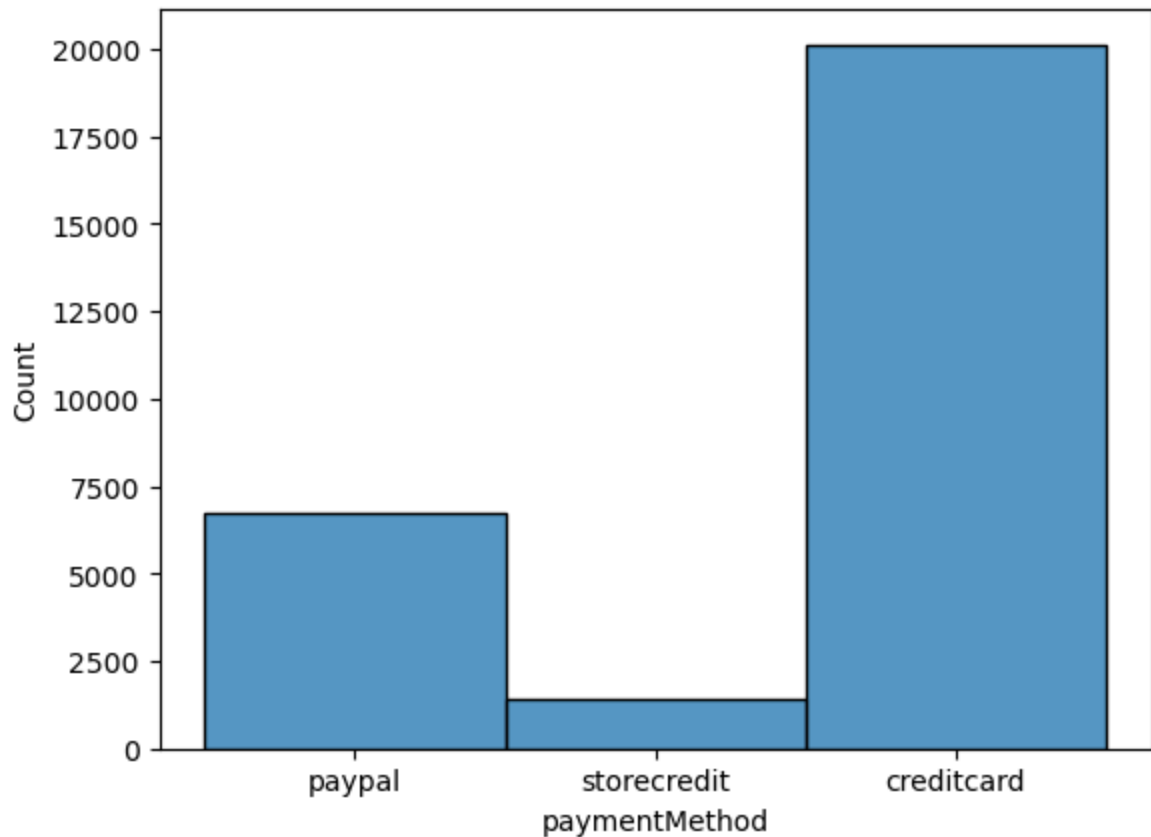
4. Exploratory Data Analysis (EDA)

- Statistical summary
- Target class distribution (fraud vs. non-fraud)
- Feature-wise distributions

```
In [13]: # Univariate
# Shows Fraud vs. Non-Fraud sample counts.
sns.countplot(data=df_cleaned, x='label', palette='Dark2', hue='label')
plt.title('Count Plot')
plt.show()
```

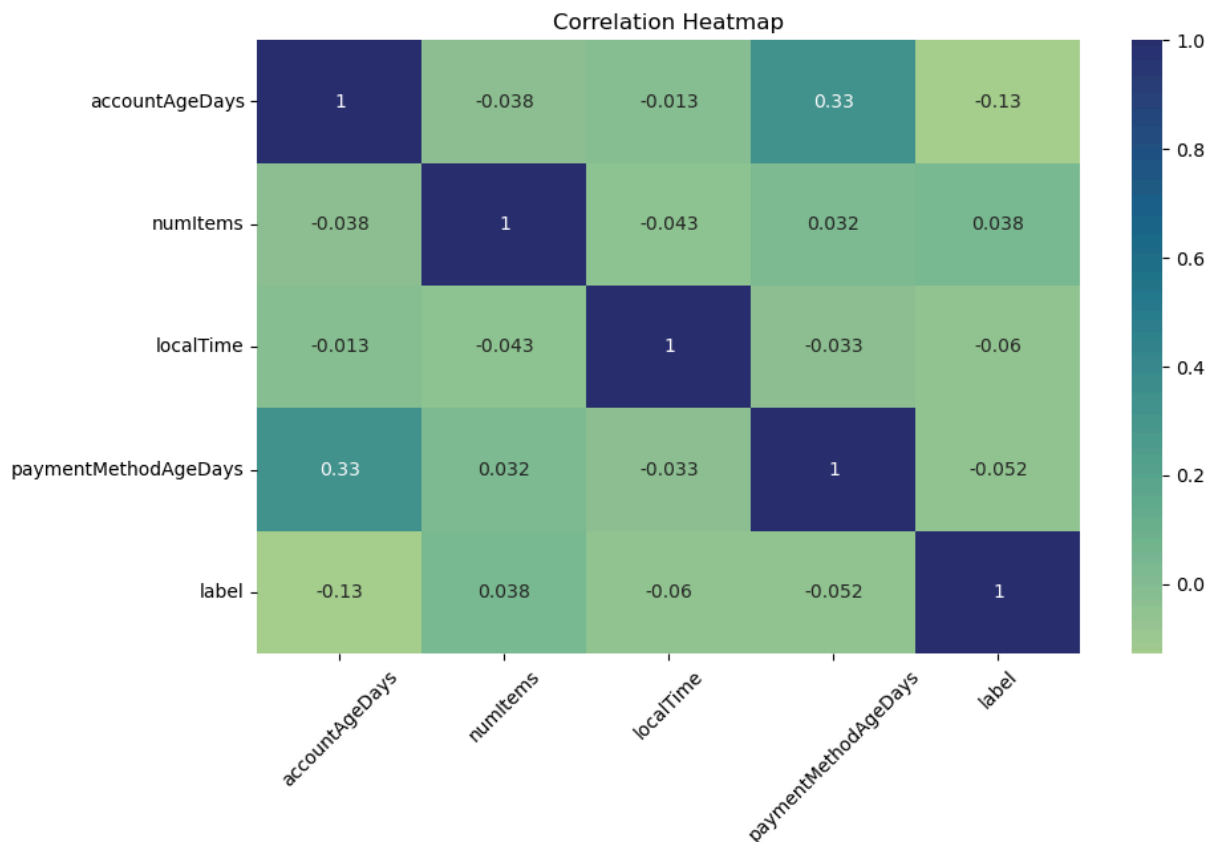


```
In [15]: # univariant
# Shows: Frequency of different payment methods.
sns.histplot(x='paymentMethod', data=df_cleaned, stat="count")
plt.show()
```



```
In [23]: # Select only numeric columns
# Shows: Correlation between numeric variables.

# here the values indicate corr between the variables +1 positive, -1 negative
numeric_df = df.select_dtypes(include='number')
correlation=numeric_df.corr()
plt.figure(figsize=(10,6))
sns.heatmap(correlation, annot=True, cmap='crest') # display the data values
plt.title('Correlation Heatmap')
plt.xticks(rotation=45)
plt.show()
```



Graphical EDA

- Boxplot of Number of Items by Fraud Status
- Boxplot of Account Age by Fraud Status and Payment Method
- Violin Plot of Payment Method Age by Fraud Status

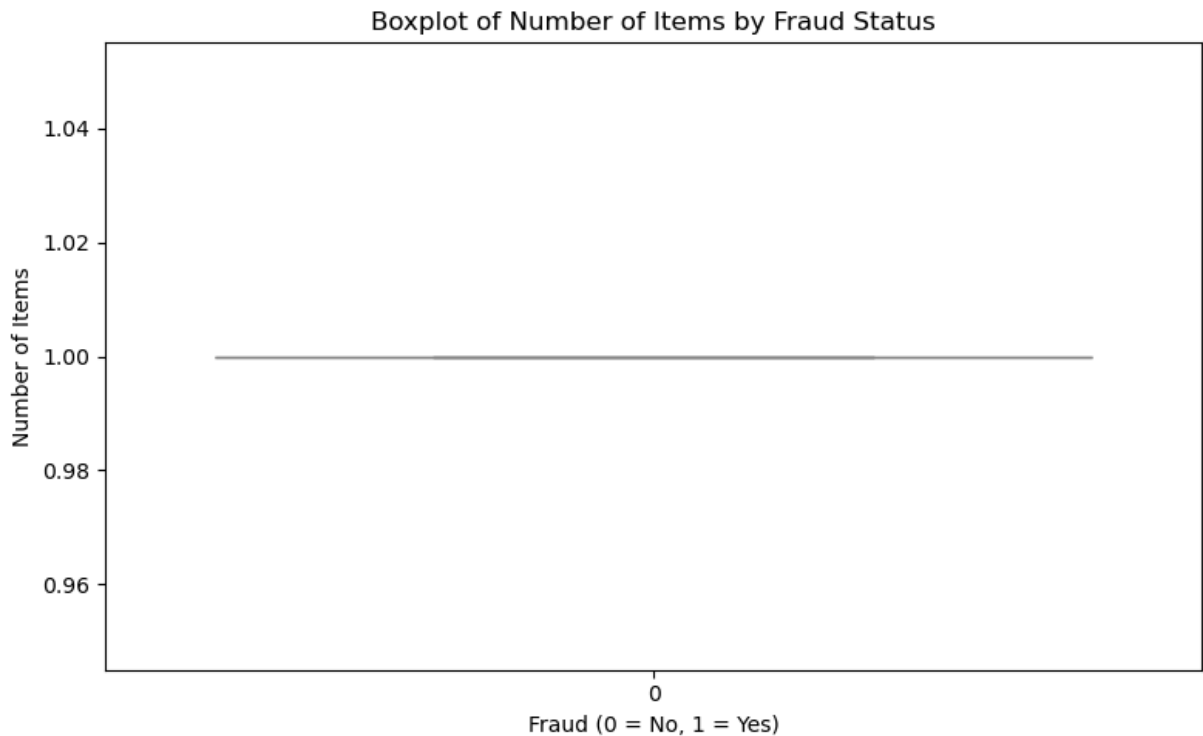
```
In [27]: # box plot Shows: Number of items per transaction based on fraud status.
plt.figure(figsize=(8, 5))
sns.boxplot(data=df_cleaned, x='label', y='numItems', palette='coolwarm')

plt.title("Boxplot of Number of Items by Fraud Status")
plt.xlabel("Fraud (0 = No, 1 = Yes)")
plt.ylabel("Number of Items")
plt.tight_layout()
plt.show()
```

C:\Users\ATHARVA\AppData\Local\Temp\ipykernel_11844\1502099584.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.


```
sns.boxplot(data=df_cleaned, x='label', y='numItems', palette='coolwarm')
<Figure size 800x500 with 0 Axes>
```



In [29]: '''After removing outliers, the number of items in both fraud and non-fraud transactions became almost the same – mostly 1.\nSo the boxplot shows a single line at 1.0 for both categories.\nThis means that, without the outliers, there's no variation in item count, and it's not a useful feature for distinguishing fraud.\nthis line is a median and represents the whole data'''

Out[29]: "After removing outliers, the number of items in both fraud and non-fraud transactions became almost the same – mostly 1.\nSo the boxplot shows a single line at 1.0 for both categories.\nThis means that, without the outliers, there's no variation in item count, and it's not a useful feature for distinguishing fraud.\nthis line is a median and represents the whole data"

In [39]: `# Boxplot of amount by fraud status`
`plt.figure(figsize=(10, 5))`
`sns.boxplot(data=df_cleaned, x='label', y='accountAgeDays', hue='paymentMethod')`
`plt.title("Boxplot of Account age (days) by Fraud Status")`
`plt.xlabel("Fraud (0 = No, 1 = Yes)")`
`plt.ylabel("Account Age(days)")`
`plt.legend(loc='upper right') # it was shown inside the graph for saving space`
`plt.show()`



In [35]: `''' This boxplot was supposed to show both fraud and non-fraud groups, but a`
`So this plot compares account age across different payment methods for only`

Out[35]: `' This boxplot was supposed to show both fraud and non-fraud groups, but af`
`ter removing outliers, only the non-fraud data remains visible.\nSo this pl`
`ot compares account age across different payment methods for only the non-f`
`raud transactions.'`

In [1]: `'''This boxplot compares the account age (in days) for fraud and non-fraud t`
`It shows whether fraud is more common in newer or older accounts, and whethe`
`The boxes are split by color using the hue parameter, which represents diffe`
`After removing outliers the fraud cases are disappeared this`
`suggest that they often involve accounts with very low account age(newer acc`


Out[1]: `'This boxplot compares the account age (in days) for fraud and non-fraud tr`
`ansactions, grouped by payment method.\nIt shows whether fraud is more comm`
`on in newer or older accounts, and whether that depends on which payment me`
`thod was used.\nThe boxes are split by color using the hue parameter, which`
`represents different payment methods like PayPal, Store Credit, and Credit`
`Card.\nAfter removing outliers the fraud cases are disappeared this \nsugge`
`st that they often involve accounts with very low account age(newer account`
`s), which were treated as outliers.'`

In [43]: `# violin plot`
`plt.figure(figsize=(8, 5))`
`sns.violinplot(data=df_cleaned, x='label', y='paymentMethodAgeDays', palette`
`plt.title("Violin plot of Payment Method Age by Fraud Status")`
`plt.xlabel("Fraud (0 = No, 1 = Yes)")`
`plt.ylabel("Payment Method Age Days")`
`plt.show()`

```
C:\Users\ATHARVA\AppData\Local\Temp\ipykernel_11844\493599876.py:4: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.violinplot(data=df_cleaned, x='label', y='paymentMethodAgeDays', palette='Set2')
```



```
In [45]: ''' This violin plot compares how long payment methods had been used in fraud
We see that most values are clustered near the bottom – around 0 to 50 days.
This means that fraud tends to happen more with newer payment methods.
The fraud side (label = 1) is almost flat or missing, showing very few fraud
```

```
Out[45]: ' This violin plot compares how long payment methods had been used in fraud
vs. non-fraud transactions.\nWe see that most values are clustered near the
bottom – around 0 to 50 days.\nThis means that fraud tends to happen more w
ith newer payment methods.\nThe fraud side (label = 1) is almost flat or mi
ssing, showing very few fraud cases left after cleaning.'
```

5. Encoding Categorical Variables

- Label Encoding: paymentMethod

```
In [47]: # label Encoding Converts categorical payment method to numeric values.
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
df_cleaned['paymentMethod'] = le.fit_transform(df_cleaned.paymentMethod.values)
df_cleaned.head()
```



```
Out[47]:
```

	accountAgeDays	numItems	localTime	paymentMethod	paymentMethodAge
0	29	1	4.745402	1	28.
1	725	1	4.742303	2	0.
2	845	1	4.921318	0	0.
3	503	1	4.886641	0	0.
4	2000	1	5.040929	0	0.

```
In [49]: df_cleaned.paymentMethod.value_counts()
```

```
Out[49]: paymentMethod
0      20126
1       6750
2       1400
Name: count, dtype: int64
```

```
In [51]: # One-hot Encoding create a column for each category

one_hot = pd.get_dummies(df_cleaned['paymentMethod'])
one_hot
```

```
Out[51]:
```

	0	1	2
0	False	True	False
1	False	False	True
2	True	False	False
3	True	False	False
4	True	False	False
...
39212	True	False	False
39213	True	False	False
39216	True	False	False
39218	True	False	False
39220	False	True	False

28276 rows × 3 columns

```
In [53]: df_cleaned.head(10)
```

Out[53]:	accountAgeDays	numItems	localTime	paymentMethod	paymentMethodAge
0	29	1	4.745402	1	28.
1	725	1	4.742303	2	0.
2	845	1	4.921318	0	0.
3	503	1	4.886641	0	0.
4	2000	1	5.040929	0	0.
5	119	1	4.962055	1	0.
6	2000	1	4.921349	1	0.
7	371	1	4.876771	0	0.
8	2000	1	4.748314	0	0.
9	4	1	4.461622	0	0.

In [31]: `pip install -U scikit-learn`

Requirement already satisfied: scikit-learn in c:\users\atharva\anaconda3\lib\site-packages (1.6.1)

Collecting scikit-learn

Using cached scikit_learn-1.7.1-cp312-cp312-win_amd64.whl.metadata (11 kB)

Requirement already satisfied: numpy>=1.22.0 in c:\users\atharva\anaconda3\lib\site-packages (from scikit-learn) (1.26.4)

Requirement already satisfied: scipy>=1.8.0 in c:\users\atharva\anaconda3\lib\site-packages (from scikit-learn) (1.13.1)

Requirement already satisfied: joblib>=1.2.0 in c:\users\atharva\anaconda3\lib\site-packages (from scikit-learn) (1.4.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\atharva\anaconda3\lib\site-packages (from scikit-learn) (3.5.0)

Using cached scikit_learn-1.7.1-cp312-cp312-win_amd64.whl (8.7 MB)

Installing collected packages: scikit-learn

Attempting uninstall: scikit-learn

Found existing installation: scikit-learn 1.6.1

Uninstalling scikit-learn-1.6.1:

Successfully uninstalled scikit-learn-1.6.1

Successfully installed scikit-learn-1.7.1

Note: you may need to restart the kernel to use updated packages.

WARNING: Failed to remove contents in a temporary directory 'C:\Users\ATHARVA\anaconda3\Lib\site-packages\~=learn'.

You can safely remove it manually.

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

sklearn-compat 0.1.3 requires scikit-learn<1.7,>=1.2, but you have scikit-learn 1.7.1 which is incompatible.

6. Feature Scaling

- Standardization
- Normalization

```
In [32]: from sklearn.preprocessing import StandardScaler  
from sklearn.preprocessing import MinMaxScaler
```

```
In [55]: df_cleaned.head()  
df_cleaned.describe().round(2)
```

```
Out[55]:
```

	accountAgeDays	numItems	localTime	paymentMethod	paymentMeth
count	28276.00	28276.0	28276.00		28276.00
mean	754.78	1.0	4.85		0.34
std	784.98	0.0	0.16		0.57
min	2.00	1.0	4.46		0.00
25%	51.00	1.0	4.75		0.00
50%	365.00	1.0	4.89		0.00
75%	1589.00	1.0	4.96		1.00
max	2000.00	1.0	5.04		2.00

```
In [57]: # normalization  
  
new_df = pd.DataFrame(df_cleaned, columns = ['paymentMethod', 'paymentMethodA  
new_df.head(5)
```

```
Out[57]:
```

	paymentMethod	paymentMethodAgeDays
0	1	28.204861
1	2	0.000000
2	0	0.000000
3	0	0.000000
4	0	0.000000

```
In [35]: scalar = MinMaxScaler() #instantiating the minmaxscalar() function  
normalized_df = scalar.fit_transform(new_df)  
print(normalized_df)
```

```
[5.00000000e-01 1.41053888e-02]
[1.00000000e+00 0.00000000e+00]
[0.00000000e+00 0.00000000e+00]
...
[0.00000000e+00 0.00000000e+00]
[0.00000000e+00 1.18066081e-01]
[5.00000000e-01 3.47295058e-07]]
```

```
In [36]: # standardization
scalar = StandardScaler() #instantiating the standardscalar() function
standardized_df = scalar.fit_transform(new_df)
print(standardized_df)
```

```
[[ 1.17536144 -0.33303221]
 [ 2.94227791 -0.43249725]
 [-0.59155504 -0.43249725]
...
 [-0.59155504 -0.43249725]
 [-0.59155504  0.40005321]
 [ 1.17536144 -0.4324948 ]]
```

```
In [37]: !pip install -U scikit-learn imbalanced-learn
```

```
Requirement already satisfied: scikit-learn in c:\users\atharva\anaconda3\li
b\site-packages (1.7.1)
Requirement already satisfied: imbalanced-learn in c:\users\atharva\anaconda
3\lib\site-packages (0.13.0)
Requirement already satisfied: numpy>=1.22.0 in c:\users\atharva\anaconda3\l
ib\site-packages (from scikit-learn) (1.26.4)
Requirement already satisfied: scipy>=1.8.0 in c:\users\atharva\anaconda3\li
b\site-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in c:\users\atharva\anaconda3\l
ib\site-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\atharva\anac
onda3\lib\site-packages (from scikit-learn) (3.5.0)
Requirement already satisfied: sklearn-compat<1,>=0.1 in c:\users\atharva\an
aconda3\lib\site-packages (from imbalanced-learn) (0.1.3)
Collecting scikit-learn
  Using cached scikit_learn-1.6.1-cp312-cp312-win_amd64.whl.metadata (15 kB)
Using cached scikit_learn-1.6.1-cp312-cp312-win_amd64.whl (11.1 MB)
Installing collected packages: scikit-learn
  Attempting uninstall: scikit-learn
    Found existing installation: scikit-learn 1.7.1
    Uninstalling scikit-learn-1.7.1:
      Successfully uninstalled scikit-learn-1.7.1
Successfully installed scikit-learn-1.6.1
```

7. Train-Test Split

- 80-20 split
- Stratified sampling to handle class imbalance

```
In [61]: # Train-Test Split
from sklearn.model_selection import train_test_split
X = df_cleaned.drop('label', axis=1)
```

```

y = df_cleaned['label']

# X_train, y_train: 80% of the data for training
# X_test, y_test: 20% of the data for testing
# test_size=0.2: Keep 20% of data for testing.
# stratify=y: Ensures the same ratio of fraud and non-fraud in both train and test sets
# random_state=42: Fixes the random split so you get the same result every time

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)

```

```

In [39]: #print(X_train.dtypes)

import sklearn
import imblearn

print("scikit-learn version:", sklearn.__version__)
print("imbalanced-learn version:", imblearn.__version__)

```

```

scikit-learn version: 1.6.1
imbalanced-learn version: 0.13.0

```

8. Handling Class Imbalance

- Undersampling (RandomUnderSampler)
- Oversampling (SMOTE)

```

In [40]: # undersampling - majority class(non fraud) is reduced to minority class(fraud)
from imblearn.under_sampling import RandomUnderSampler

undersample = RandomUnderSampler(random_state=42)
X_under, y_under = undersample.fit_resample(X_train, y_train)
# fit_resample() does two things: Fit the resampling logic (e.g., how many samples to remove)
# Resample the dataset accordingly and return the balanced X and y

print("Before:", np.bincount(y_train)) # counts the number of occurrences of each class
print("After:", np.bincount(y_under)) # y_under is the resampled version of y_train

```

```

Before: [30928  448]
After: [448 448]

```

```

In [69]: '''We want to reduce the majority class, we need to count labels, not features.
Features (X): These are the input variables used to make a prediction.
Example: account age, number of items, payment method, etc.

Label (y): This is the output variable (what we want to predict).
Example: whether a transaction is fraud or not (0 or 1).'''

```

```

Out[69]: 'We want to reduce the majority class, we need to count labels, not features.
Features (X): These are the input variables used to make a prediction.
Example: account age, number of items, payment method, etc.

Label (y): This is the output variable (what we want to predict).
Example: whether a transaction is fraud or not (0 or 1).'

```

```

In [41]: # Oversampling - Increases the number of samples in the minority class (fraud)
from imblearn.over_sampling import SMOTE # Synthetic Minority Over-sampling

```

```

# It creates fake but realistic examples of the minority class.
# Instead of copying the same rows, SMOTE looks at a point's nearest neighbors
smote = SMOTE(random_state=42)
X_smote, y_smote = smote.fit_resample(X_train, y_train)

print("Before:", np.bincount(y_train))
print("After:", np.bincount(y_smote))

```

Before: [30928 448]

After: [30928 30928]

```

In [62]: from sklearn.metrics import classification_report, confusion_matrix
# classifies data accuracy according to the models used

```

```

# from sklearn.datasets import make_regression # linear regression

```

```

In [71]: '''We scale data when using distance-based or gradient-based models like KNN,
Tree-based models like Random Forest don't need scaling because they split b

```

```

Out[71]: 'We scale data when using distance-based or gradient-based models like KNN,
SVM, or logistic regression.\nTree-based models like Random Forest don't ne
ed scaling because they split based on thresholds, not distances.'

```

Model Building & Evaluation

Model 1: Random Forest Classifier

- Training
- Confusion Matrix
- Performance Metrics

```

In [57]: # Random Forest

```

```

from sklearn.ensemble import RandomForestClassifier
model_rf = RandomForestClassifier() # Train Random Forest model
model_rf.fit(X_smote, y_smote) # or use X_under, y_under

# Predict
y_pred_rf = model_rf.predict(X_test)

# Classification Report
from sklearn.metrics import classification_report
print("Random Forest:\n", classification_report(y_test, y_pred_rf))

```

Random Forest:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7733
1	1.00	1.00	1.00	112
accuracy			1.00	7845
macro avg	1.00	1.00	1.00	7845
weighted avg	1.00	1.00	1.00	7845

Model 2: Support Vector Machine (SVM)

- Training
- Confusion Matrix
- Performance Metrics

```
In [60]: from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler

# Scale training data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_smote) # or X_under
X_test_scaled = scaler.transform(X_test)

# Train SVM model
model_svm = SVC()
model_svm.fit(X_train_scaled, y_smote) # or y_under

# Predict
y_pred_svm = model_svm.predict(X_test_scaled)

# Classification Report
from sklearn.metrics import classification_report
print("SVM:\n", classification_report(y_test, y_pred_svm))
```

SVM:

	precision	recall	f1-score	support
0	1.00	0.86	0.93	7733
1	0.09	1.00	0.17	112
accuracy			0.86	7845
macro avg	0.55	0.93	0.55	7845
weighted avg	0.99	0.86	0.92	7845

Model 3: K-Nearest Neighbors (KNN)

- Training
- Confusion Matrix
- Performance Metrics

```
In [62]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler

# Scale training data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_smote) # or X_under
X_test_scaled = scaler.transform(X_test)

# Train KNN model
model_knn = KNeighborsClassifier()
model_knn.fit(X_train_scaled, y_smote) # or y_under

# Predict
y_pred_knn = model_knn.predict(X_test_scaled)

# Classification Report
from sklearn.metrics import classification_report
print("KNN:\n", classification_report(y_test, y_pred_knn))

# 0 is class 0 (non-fraud)
# 1 is class 1 (fraud)

# accuracy = overall correct predictions
# macro avg = average across classes (treat all classes equally)
# weighted avg = average considering how many samples are in each class
```

KNN:

	precision	recall	f1-score	support
0	1.00	0.99	1.00	7733
1	0.65	0.95	0.77	112
accuracy			0.99	7845
macro avg	0.83	0.97	0.88	7845
weighted avg	0.99	0.99	0.99	7845

Precision: How precise your fraud predictions are

Recall: How many real frauds you caught

F1-score: Balance between precision and recall

Accuracy: Overall performance

Plotting confusion matrix

```
In [64]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```



```


# For KNN
cm_knn = confusion_matrix(y_test, y_pred_knn)
disp_knn = ConfusionMatrixDisplay(confusion_matrix=cm_knn)
disp_knn.plot(cmap='Blues')
plt.title("KNN Confusion Matrix")
plt.show()



# For SVM
cm_svm = confusion_matrix(y_test, y_pred_svm)
disp_svm = ConfusionMatrixDisplay(confusion_matrix=cm_svm)
disp_svm.plot(cmap='Purples')
plt.title("SVM Confusion Matrix")
plt.show()

# For Random Forest
cm_rf = confusion_matrix(y_test, y_pred_rf)
disp_rf = ConfusionMatrixDisplay(confusion_matrix=cm_rf)
disp_rf.plot(cmap='Greens')
plt.title("Random Forest Confusion Matrix")
plt.show()

# the graph below tells that how many values are true positive and true negative

```





```
In [79]: # TN    True Negatives (non-fraud correctly predicted)
# TP    True Positives (fraud correctly predicted)
# FP    False Positives (non-fraud wrongly marked as fraud)
# FN    False Negatives (fraud missed as non-fraud)

#
#           Predicted No (0)      Predicted Yes (1)

# Actual No (0)   True Negative (TN)   False Positive (FP)
# Actual Yes (1)  False Negative (FN)  True Positive (TP)

# Based on the confusion matrices, KNN is the most balanced model – it detects
# SVM over-predicts fraud, and Random Forest misses all frauds."
```

Model Graphs

Model Comparison & Metrics Visualization

- Accuracy
- Precision
- Recall
- F1 Score

```
In [86]: # Metric---> Type of performance measurement (Precision, Recall, Accuracy, etc)
# Score---> The actual result or value for that metric (e.g., 0.89 = 89%)

from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Metrics for each model (using predictions on original y_test)
metrics = {
    'KNN': {
        'Precision': precision_score(y_test, y_pred_knn),
        'Recall': recall_score(y_test, y_pred_knn),
        'F1-Score': f1_score(y_test, y_pred_knn),
        'Accuracy': accuracy_score(y_test, y_pred_knn)
    },
    'SVM': {
        'Precision': precision_score(y_test, y_pred_svm),
        'Recall': recall_score(y_test, y_pred_svm),
        'F1-Score': f1_score(y_test, y_pred_svm),
        'Accuracy': accuracy_score(y_test, y_pred_svm)
    },
    'Random Forest': {
        'Precision': precision_score(y_test, y_pred_rf),
        'Recall': recall_score(y_test, y_pred_rf),
        'F1-Score': f1_score(y_test, y_pred_rf),
        'Accuracy': accuracy_score(y_test, y_pred_rf)
    }
}
```

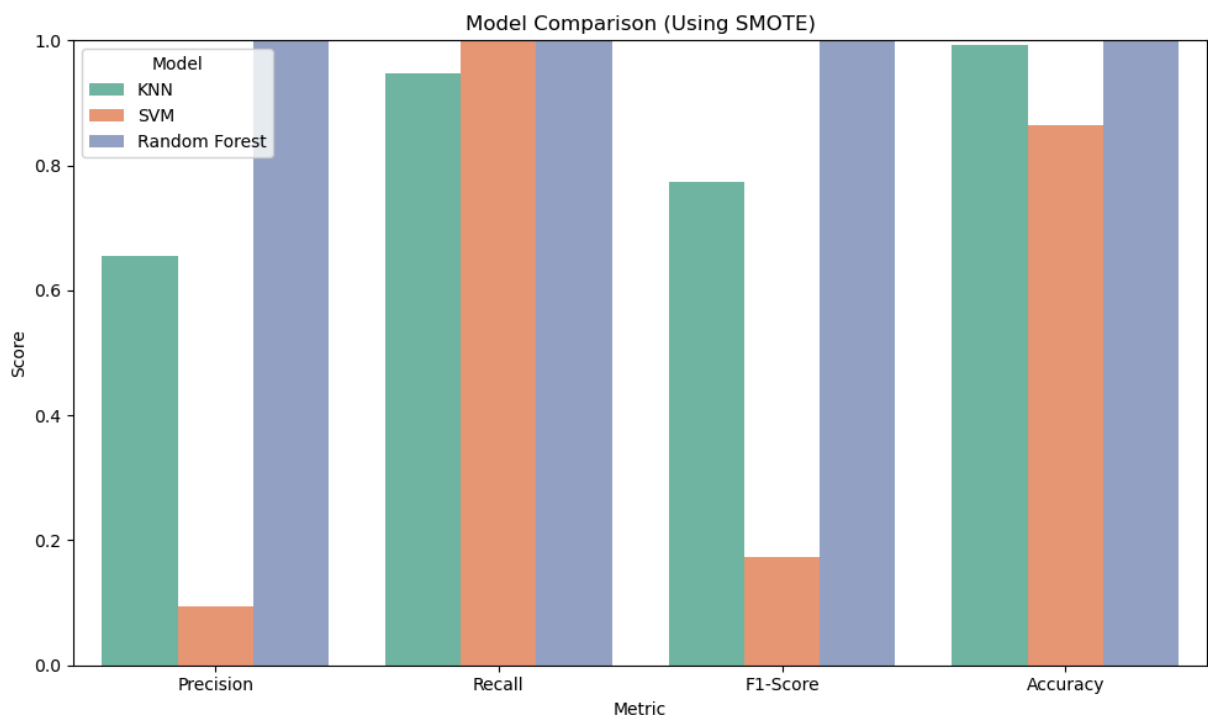
```

# Convert to DataFrame
# columns are models (KNN, SVM, etc.) ; rows are metric names
# .T transposes the DataFrame
# After transpose, the row labels (KNN, SVM, etc.) are in the index reset_index
# Melts (unpivots) the DataFrame from wide to long format ; Keeps the 'index'
# unpivot column metrics to row metrics

df_metrics = pd.DataFrame(metrics).T.reset_index().melt(id_vars='index')
df_metrics.columns = ['Model', 'Metric', 'Score'] # Renames the columns from

# Plot
plt.figure(figsize=(10, 6))
sns.barplot(data=df_metrics, x='Metric', y='Score', hue='Model', palette='Set1')
plt.title("Model Comparison (Using SMOTE)")
plt.ylim(0, 1)
plt.ylabel("Score")
plt.legend(title='Model')
plt.tight_layout()
plt.show()

```



```

In [88]: # This helps you visually evaluate which model performs best on which metric
# For fraud detection, we focus more on F1-Score and Recall because it's important
# So, the model with the highest bars in those metrics would be the most effective

# Accuracy--> Overall correct predictions
# Precision--> Of all fraud predictions, how many were right
# Recall--> Of all actual frauds, how many were caught
# F1-Score--> Overall fraud detection effectiveness

```

In []:

